

## NEURAL NETWORKS ASSIGNMENT-3

### Deep Learning Image Classification with CNN

Name: Saharsha Muddagauni STD ID: 700742790

GITHUB LINK: [https://github.com/sxm26790/ASSIGNMENT3\\_NEURAL](https://github.com/sxm26790/ASSIGNMENT3_NEURAL)

To perform image classification with a CNN, we'll use the Keras library, which provides an easy-to-use API for building and training deep learning models. Let's go through the implementation step by step:

Import the required libraries:

```
import numpy as np
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.constraints import maxnorm
from keras.optimizers import SGD
from keras.layers.convolutional import Conv2D, MaxPooling2D
from keras.utils import np_utils
```

```
np.random.seed(7)
```

Load and preprocess the data: We'll use the CIFAR-10 dataset, which contains 50,000 training images and 10,000 test images of 10 different classes.

```
(X_train, y_train), (X_test, y_test) = cifar10.load_data()
```

```
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 [=====] - 3s 0us/step
```

```
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0
```

```
y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)
num_classes = y_test.shape[1]
```

Build the CNN model:

```
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2), padding='same'))
model.add(Flatten())
model.add(Dense(512, activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))
```

```
sgd = SGD(learning_rate=0.01, momentum=0.9, decay=1e-6)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 32, 32, 32)	896
dropout (Dropout)	(None, 32, 32, 32)	0
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
flatten (Flatten)	(None, 8192)	0
dense (Dense)	(None, 512)	4194816
dropout_1 (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 10)	5130

=====

Total params: 4,210,090  
Trainable params: 4,210,090  
Non-trainable params: 0

---

None

```
epochs = 5
batch_size = 32
model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch_size)
```

```
Epoch 1/5
1563/1563 [=====] - 164s 105ms/step - loss: 1.7276 - accuracy: 0.3737 - val_loss: 1.4764 - val_accuracy: 0.4738
Epoch 2/5
1563/1563 [=====] - 168s 108ms/step - loss: 1.4083 - accuracy: 0.4939 - val_loss: 1.2520 - val_accuracy: 0.5621
Epoch 3/5
1563/1563 [=====] - 168s 108ms/step - loss: 1.2226 - accuracy: 0.5638 - val_loss: 1.1403 - val_accuracy: 0.5962
Epoch 4/5
1563/1563 [=====] - 167s 107ms/step - loss: 1.0943 - accuracy: 0.6109 - val_loss: 1.0451 - val_accuracy: 0.6321
Epoch 5/5
1563/1563 [=====] - 169s 108ms/step - loss: 0.9879 - accuracy: 0.6515 - val_loss: 1.1353 - val_accuracy: 0.6127
<keras.callbacks.History at 0x781917f1de40>
```

```
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))
```

Accuracy: 61.27%

```
import numpy as np
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers.convolutional import Conv2D, MaxPooling2D
from keras.constraints import maxnorm
from keras.utils import np_utils
from keras.optimizers import SGD
```

```
# Fix random seed for reproducibility
np.random.seed(7)
```

```
# Load data
(X_train, y_train), (X_test, y_test) = cifar10.load_data()
```

```
# Normalize inputs from 0-255 to 0.0-1.0
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0
```

```
# One hot encode outputs
y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)
num_classes = y_test.shape[1]
```

```
# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))
```

```
# Compile model
epochs = 5
learning_rate = 0.01
decay_rate = learning_rate / epochs
sgd = SGD(lr=learning_rate, momentum=0.9, decay=decay_rate, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())
```

```
# Fit the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)
```

```
# Evaluate the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1] * 100))
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_2 (Conv2D)	(None, 32, 32, 32)	896
dropout_2 (Dropout)	(None, 32, 32, 32)	0
conv2d_3 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d_1 (MaxPooling 2D)	(None, 16, 16, 32)	0
conv2d_4 (Conv2D)	(None, 16, 16, 64)	18496
dropout_3 (Dropout)	(None, 16, 16, 64)	0
conv2d_5 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_2 (MaxPooling 2D)	(None, 8, 8, 64)	0
conv2d_6 (Conv2D)	(None, 8, 8, 128)	73856
dropout_4 (Dropout)	(None, 8, 8, 128)	0
conv2d_7 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_3 (MaxPooling 2D)	(None, 4, 4, 128)	0
flatten_1 (Flatten)	(None, 2048)	0
dropout_5 (Dropout)	(None, 2048)	0
dense_2 (Dense)	(None, 1024)	2098176
dropout_6 (Dropout)	(None, 1024)	0
dense_3 (Dense)	(None, 512)	524800
dropout_7 (Dropout)	(None, 512)	0
dense_4 (Dense)	(None, 10)	5130

=====

Total params: 2,915,114  
Trainable params: 2,915,114  
Non-trainable params: 0

/usr/local/lib/python3.10/dist-packages/keras/optimizers/legacy/gradient\_descent.py:114: UserWarning: The `lr` argument is deprecated, use  
super().\_\_init\_\_(name, \*\*kwargs)

None  
Epoch 1/5  
1563/1563 [=====] - 342s 218ms/step - loss: 1.9313 - accuracy: 0.2843 - val\_loss: 1.7142 - val\_accuracy: 0.3886  
Epoch 2/5  
1563/1563 [=====] - 345s 221ms/step - loss: 1.5361 - accuracy: 0.4434 - val\_loss: 1.4258 - val\_accuracy: 0.4928  
Epoch 3/5

```
import numpy
# Predict the first 4 images of the test data
predictions = model.predict(X_test[:4])
# Convert the predictions to class labels
predicted_labels = numpy.argmax(predictions, axis=1)
# Convert the actual labels to class labels
actual_labels = numpy.argmax(y_test[:4], axis=1)

# Print the predicted and actual labels for the first 4 images
print("Predicted labels:", predicted_labels)
print("Actual labels: ", actual_labels)
```

1/1 [=====] - 0s 202ms/step  
Predicted labels: [3 1 8 0]  
Actual labels: [3 8 8 0]

```
import matplotlib.pyplot as plt

# Plot the training and validation loss
```

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'val'], loc='upper right')
plt.show()
```

```
# Plot the training and validation accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['train', 'val'], loc='lower right')
plt.show()
```



